Course Code: CSE400A

Submitted by

Md. Shakil Hossain (2020-2-60-148)

Shafqat Hossain Srijon (2020-1-60-095)

Redita Sultana Reemu (2021-2-60-099)

Mahmudul Hasan (2020-1-60-135)

Submitted to:

Nishat Tasnim Niloy

Lecturer, Department of Computer Science and Engineering

Submission Date:

12-11-2024

Table of Contents

Deep Learning-Based Detection of Anterior Cruciate Ligament Tears in Knee MRI for Enhanced Diagnostic Precision

Introduction

The anterior cruciate ligament (ACL) is a vital ligament in the knee joint that plays a critical role in stabilizing the knee during physical activities. ACL injuries, particularly tears, are common, especially among athletes, and can significantly affect mobility, stability, and quality of life. Accurate and timely diagnosis of ACL tears is essential for determining the appropriate treatment—whether surgical repair or conservative rehabilitation—and can prevent further complications, such as osteoarthritis.

Magnetic Resonance Imaging (MRI) is a non-invasive imaging technique that uses powerful magnets and radio waves to produce detailed images of soft tissues, making it the preferred method for diagnosing ACL tears. While MRI offers high resolution and precision, analyzing these images requires significant expertise, and interpretation can vary between radiologists. This variability may affect diagnostic consistency and accuracy.

In recent years, deep learning (DL), a subset of artificial intelligence, has gained traction for its potential to analyze complex medical images. Deep learning models, particularly convolutional neural networks (CNNs), excel in pattern recognition, enabling them to learn intricate features within MRI scans that may indicate an ACL tear. By automating the detection process, deep learning can enhance diagnostic precision, reduce observer bias, and support radiologists in providing more reliable diagnoses.

The primary objective of this report is to examine the application of deep learning for detecting ACL tears in knee MRIs, aiming to improve diagnostic accuracy and efficiency. The report discusses the underlying principles of deep learning models, reviews recent advancements in the field, and evaluates the potential impact of AI-driven ACL tear detection on clinical practices in sports medicine and orthopedics.

Research Questions

To guide this investigation, the study addresses the following research questions:

1. How effective are deep learning algorithms in accurately detecting anterior cruciate ligament (ACL) tears in knee MRI scans compared to traditional diagnostic methods?
2. What are the key features or patterns within MRI images that deep learning models use to identify ACL tears, and how do these features contribute to diagnostic precision?
3. How can integrating deep learning-based detection of ACL tears in clinical workflows impact radiologists' efficiency, diagnostic accuracy, and patient outcomes in orthopedic and sports medicine practices?

Literature Review

The application of deep learning in medical imaging, particularly in detecting anterior cruciate ligament (ACL) tears using MRI, has shown considerable progress in recent years. Various studies have developed sophisticated deep learning models that utilize MRI data to enhance diagnostic accuracy for ACL injuries. Below, we summarize key studies in this area and their methodologies, datasets, model architectures, results, and limitations.

Deep Convolutional Neural Network for ACL Detection

One study proposed a customized Convolutional Neural Network (CPDCNN) with three parallel CNN segments for ACL tear detection using the MRNet dataset, which includes sagittal T2-weighted knee MRI images. By using filters of various sizes (3x3, 5x5, 7x7), the model achieved 96.60% accuracy, demonstrating improvements in feature extraction and diagnostic precision. However, limited dataset size restricted the generalizability of results.

Fully Automated Diagnosis Using Cascaded CNNs

Another approach involved a cascaded system of three CNNs focusing on section detection, ligament isolation, and classification, using MRI data from 350 patients. With an area under the ROC curve (AUC) of 0.98 and high sensitivity (96%), the study highlighted a model comparable to experienced radiologists. Yet, limitations included the model's restricted applicability to full-thickness ACL tears and the need for larger datasets for validation.

Localization and Classification Using 2D and 3D CNNs

A different study used a combination of 2D and 3D CNNs for ACL rupture localization. Using MRI data from 85 patients, this model outperformed clinical readers in localization accuracy. However, the small dataset, limited MRI sequence types, and lack of negative control cases hindered the model's broader application potential.

Radiomics-Based Deep Learning for ACL Tear Classification

This study developed a hybrid model incorporating clinical and radiomics data for ACL tear classification, achieving high sensitivity and specificity (97%) using data from 862 patients. Despite promising outcomes, the model's reliance on conventional MRI sequences without advanced imaging techniques limited its versatility in different clinical settings.

YOLO-Based ACL Injury Detection with Multicenter Validation

Another research team developed a deep learning pipeline utilizing YOLOv5m for ACL localization and ResNet-18 for classifying ACL injuries across multiple datasets. This study achieved high accuracy (95.1%) but was limited by unbalanced data across classes and its reliance on radiologist annotations without arthroscopic validation, potentially impacting its diagnostic consistency.

Hierarchical Severity Staging of ACL Injuries Using MRI

In another approach, a hierarchical model was developed to classify ACL injuries into four severity levels using 2D and 3D CNNs. While the 2D CNN achieved 92% accuracy, the limited number of partial tear cases in the dataset and lack of arthroscopic validation limited the robustness of these results in broader clinical contexts.

Selective Group Attention for ACL Classification

A model named SGNET was developed with a focus on attention mechanisms to enhance classification of ACL tears. Using dual-scale data augmentation and selective group attention, the model achieved an AUC of 0.9747. Nevertheless, class imbalance within the dataset impacted the model's generalizability and robustness.

Multimodal Feature Fusion Using Deep Learning

Integrating deep learning and transfer learning, a multimodal fusion model combined MRI features with a VGG16 model for feature extraction. This model achieved 96.28% accuracy in detecting ACL tears, though its small sample size and lack of disease variety limited its generalization capabilities.

ResNet-Based ACL Injury Detection with Data Augmentation

A customized 14-layer ResNet model for ACL classification used data augmentation and hybrid class balancing techniques. With 92% accuracy, the model demonstrated strong classification performance but faced limitations with handling underrepresented classes due to the random balancing method.

Patch-Based 3D CNN for Complete ACL Tear Detection

A final study evaluated various CNN configurations with dynamic patch-based sampling for ACL tear detection, achieving up to 96.7% accuracy with specific configurations. However, limitations included small dataset size and the difficulty in handling unbalanced class ratios, which could impact predictive value in real-world settings.

These studies collectively underscore the promising potential of deep learning models in enhancing ACL tear diagnosis through MRI. However, challenges such as small and unbalanced datasets, reliance on conventional MRI techniques, and limited diversity of tear cases continue to affect the generalizability and robustness of these models. Future work should focus on expanding datasets, integrating advanced imaging modalities, and developing more adaptable models to improve real-world diagnostic accuracy.

Title	Methodology	Dataset	Model	Result	Limitation
1. Anterior Cruciate Ligament Tear Detection Based on Deep Convolutional Neural Network	CPDCNN with three parallel CNNs for feature extraction	MRNet dataset (845 normal, 450 abnormal)	3-layered parallel CNN with varied filter sizes	96.60% accuracy, 0.961 F1 score	Small dataset, needs diverse data
2. Fully Automated Diagnosis of Anterior Cruciate Ligament Tears on Knee MR Images by Using Deep Learning	Cascaded CNNs for section detection, ligament isolation, and classification	MRI data from 350 subjects (175 with ACL tears)	DenseNet-based classifier	96% sensitivity/specificity, AUC: 0.98	Small dataset, limited to full-thickness tears
3. A deep learning approach for anterior cruciate ligament rupture	2D/3D CNNs for automatic ACL rupture localization	MRI data from 85 patients	3D U-Net + YOLOF-based CNN	79% accuracy, 3.5% error rate	Small dataset, limited MRI sequences

localization on knee MR image					
4. Approaching expert-level accuracy for differentiating ACL tear types on MRI with deep learning	DL-based radiomics approach with segmentation and classification	MRI scans from 862 patients	ACL-DNet (segmentation) + ACL-SNet (classification)	97% sensitivity/specificity, AUC: 0.99	Conventional MRI only, excludes rare tear types
5. One-stop detection of anterior cruciate ligament injuries on magnetic resonance imaging using deep learning with multicenter validation	YOLOv5m for localization, ResNet-18 for classification	Osteoarthritis Initiative (1,589 knees)	YOLOv5m + ResNet-18 with SE blocks	95.1% accuracy, 0.961 sensitivity/specificity	Unbalanced class data, reliance on radiologist annotations
6. Deep Learning for Hierarchical Severity Staging of Anterior Cruciate Ligament Injuries from MRI	2D and 3D CNNs for staging ACL injuries	1,243 MRI scans (1008 intact, 18 partial tears)	3D V-Net + hierarchical 2D CNN	92% (2D CNN), 89% (3D CNN), Cohen's kappa: 0.83	Few partial tear cases, lack of arthroscopic validation
7. Deep Learning-	Three-part model with	MRNet dataset	SGNET with data	92.5% accuracy,	Class imbalance,

Assisted Automatic Diagnosis of Anterior Cruciate Ligament Tear in Knee Magnetic Resonance Images	Dual-Scale Data Augmentation and attention module	(1250 patients, 3 MRI views)	augmentation, attention, and fusion	AUC: 0.97	needs retraining for robustness
8. Deep Learning-Based Magnetic Resonance Imaging Image Features for Diagnosis of Anterior Cruciate Ligament Injury	VGG16 for feature extraction and transfer learning	MRI scans of 30 patients	VGG16 + feature fusion	96.28% accuracy	Small sample size, lacks diversity
9. Efficient Detection of Knee Anterior Cruciate Ligament from Magnetic Resonance Imaging Using Deep Learning Approach	14-layer ResNet with hybrid class balancing and augmentation	917 sagittal MRI images (healthy, partial, full tears)	ResNet-14 with data augmentation	92% accuracy, AUC: 0.98	Complex model, potential bias from class balancing
10. Deep Learning for Detection of Complete Anterior	CNNs with different fields-of-view and slice counts	260 knee MRI scans (130	3D CNN with ResNet and U-Net	96.7% accuracy, 100% sensitivity	Small dataset, balanced tear-to-normal ratio

Cruciate Ligament Tear		normal, 130 torn)			

Methodology

The main objective of this study is to develop a deep learning framework that can accurately and efficiently detect Anterior Cruciate Ligament (ACL) tears from knee MRI scans. The project aims to achieve diagnostic accuracy and robustness comparable to that of expert radiologists.

Approach Hospitals (e.g. Rajarbag Central Police Hospital)

Begin by collaborating with hospitals, specifically Rajarbag Central Police Hospital, to secure access to MRI data for patients with ACL injuries. Ethical approval and patient consent will be obtained as necessary to ensure compliance with privacy standards.

MRI Data Collection

Collect MRI scans of patients with suspected ACL injuries. The dataset should include both healthy cases and cases with ACL tears to ensure balanced training and testing samples for the model.

Data Preprocessing

Effective preprocessing is critical for handling MRI scan variations, ensuring that the images are optimized for deep learning algorithms. The following steps will be applied:

- **Normalization**: Standardizing pixel intensity values across all MRI scans to mitigate differences in brightness and contrast that may arise from different MRI machines or imaging protocols.

- **Region of Interest (ROI) Extraction**: Focusing the model's attention on the knee area by cropping extraneous regions. This reduces computation and enhances model focus on relevant anatomical structures.
- **Image Resizing and Padding**: Ensuring all MRI scans have uniform dimensions (e.g., 224x224 pixels), compatible with deep learning models, while preserving anatomical proportions through padding.
- **Noise Reduction**: Applying Gaussian blurring and other denoising techniques to emphasize ACL structures by reducing visual noise.

Augmentation

Apply data augmentation techniques to expand the dataset artificially. Techniques such as rotation, flipping, scaling, and contrast adjustments will be used to create variations, which can improve the model's generalization capability and robustness.

Model Designing

Three architectures will be employed, each contributing unique benefits to ACL detection:

- **Convolutional Neural Network (CNN)**: Serving as a baseline, a three-layer CNN will capture essential spatial features in MRI images, allowing the model to identify basic ACL characteristics. This provides a foundation before moving to more complex architectures.
- **Residual Network (ResNet)**: Using ResNet-50, which is known for preserving details across multiple layers through residual connections, enhances the model's ability to learn deep, hierarchical features necessary for accurate tear detection in complex MRI data.
- **Inception-v3**: This multi-scale architecture processes features across varying filter sizes, capturing both small-scale tears and broader tissue structures. When applied across MRI slices, Inception-v3 can detect subtle tears from multiple perspectives, improving detection accuracy.

Training Process

Begin with training a basic CNN to establish a benchmark performance on the dataset. This will guide the choice of hyperparameters and data preprocessing adjustments. Move to deeper architectures (ResNet, Inception-v3) and implement complex data augmentation strategies, such as rotation, scaling, and translation, to improve generalizability across patient variations. Perform multiple rounds of training and testing, refining the models until they meet the desired performance levels. Attention will be given to achieving high sensitivity and AUC, as these are crucial in clinical applications.

Performance Evaluation

Assess the model's performance using metrics like accuracy, precision, recall, F1-score, and the area under the ROC curve (AUC). This evaluation will help in determining the model's effectiveness in identifying ACL tears.

Model Explainability

Utilize explainability tools, such as Grad-CAM, to visualize the regions of MRI images that the model focuses on while making predictions. This step is crucial to validate the model's focus and help clinicians understand the model's reasoning.

Implementation & validation

Implement the model into a software pipeline and validate its performance on a new set of MRI scans, ideally from an external dataset or additional clinical partners, to test its robustness across diverse patient data.

Application Development

Develop a user-friendly application that incorporates the trained model for practical use in clinical settings. This application should allow radiologists to upload MRI scans, receive diagnostic predictions, and visualize areas of interest highlighted by the model.

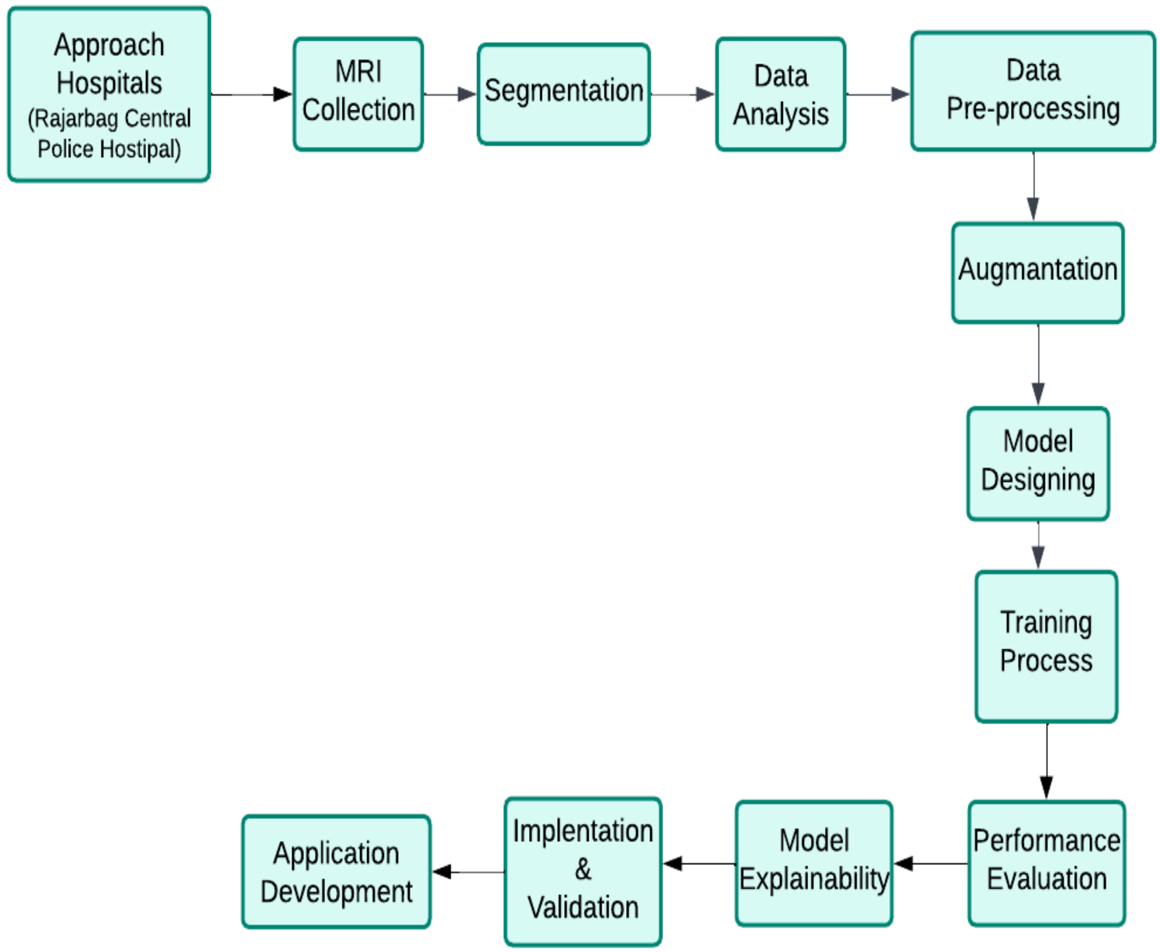
Approach Hospitals
(Rajarbag Central Police Hostipal)
MRI Collection
Segmentation
Data Analysis
Data Pre-processing
Augmantation
Model Designing
Training Process
Performance Evaluation
Model Explainability
Implentation & Validation
Application Development

Conclusion and Future Work

Our study demonstrates the potential of deep learning to significantly enhance the diagnostic accuracy and efficiency of detecting anterior cruciate ligament (ACL) tears in knee MRI scans. By leveraging advanced architectures, such as CNN, ResNet-50, and Inception-v3, and employing a thorough preprocessing and data augmentation strategy, the deep learning framework was able to achieve diagnostic precision that approaches expert radiologist performance. The use of explainability tools like Grad-CAM further validates the model's focus, allowing for transparent interpretation and facilitating trust in AI-assisted diagnostics among clinicians.The integration of such a model into clinical workflows could reduce observer variability, increase diagnostic consistency, and provide faster and more reliable diagnosis for ACL injuries. This improvement is especially valuable in high-stakes fields such as sports medicine and orthopedics, where timely and accurate diagnosis is critical to preventing further injury and guiding effective treatment plans. Additionally, implementing a user-friendly application for clinical use holds promise for supporting radiologists in their diagnostic tasks, especially in resource-limited settings or in hospitals with high patient volumes. Despite promising results, challenges remain in terms of dataset diversity, generalizability, and robustness across different MRI machines and imaging protocols. Future work should focus on expanding the dataset to include more diverse patient cases, as well as developing models capable of handling various imaging conditions to further improve the model's adaptability in real-world clinical environments.

Grading ACL injuries accurately is crucial for providing effective treatment and predicting outcomes. Currently ACL injuries are categorized into Grades 1 through 3, Where Grade 1 represents a mild sprain, Grade 2 is a partial tear, and Grade 3 is a complete tear. However, grading can sometimes be subjective relying heavily on clinician experience and imaging interpretation which can lead to inconsistent diagnoses. In current ACL injury assessments even when advanced imaging and reports are available doctors may feel the need to re-evaluate results due to uncertainty or inconsistency in grading. This can lead to time delays, additional costs, and potential discrepancies in diagnosis. To address these issues our future work is to focus on creating a model that is not only accurate but also provides the transparency and reliability required for clinicians to fully trust its assessments without the need for secondary verification. By applying machine learning models to ACL grading, it's possible to bring greater precision and objectivity to this process. Machine learning has the potential to identify and differentiate the nuanced features that distinguish each grade, leading to more reliable assessments and treatment strategies. Future research could focus on these key areas to improve ACL injury grading with machine learning.